Great Place To Work. Certified FEB 2023 – FEB 2024 INDIA

UPES

UNIVERSITY OF TOMORROW

nirf RANKING 2023

#RANKED 52 IN INDIA

#University Category

QS WORLD UNIVERSITY RANKINGS

NO.1 PVT. UNIVERSITY IN ACADEMIC REPUTATION IN INDIA

NAAC

ACCREDITED GRADE 'A' BY NAAC

PERFECT SCORE OF 150/150 AS A TESTAMENT TO EXECEPTIONAL E-LEARNING METHODS

UPES UNIVERSITY OF TOMORROW

Elements of AIML

Unit 04 – Lecture 03: Clustering, Reinforcement Learning & Semi-Supervised ML

Tofik Ali

School of Computer Science, UPES Dehradun

February 28, 2026

Topic focus: **unsupervised** and **interaction-based** learning paradigms

Repository: `https://github.com/tali7c/Elements-of-AIML`

Sources: Mueller and Massaron, *Machine Learning for Dummies*, For Dummies, 2016.;James, Witten, Hastie, and Tibshirani, *An Introduction to Statistical Learning* (2e), selected chapters.;Scikit-learn documentation, accessed 2026-02-28.

Quick Links

Clustering k-means RL Basics Semi-Supervised Summary

Agenda

Clustering

k-means

Reinforcement Learning

Semi-supervised learning

Summary

Learning Outcomes

By the end of this lecture, you should be able to:

- Explain clustering and where it is used.

Learning Outcomes

By the end of this lecture, you should be able to:

- Explain clustering and where it is used.
- Describe the k-means algorithm and its objective.

Learning Outcomes

By the end of this lecture, you should be able to:

- Explain clustering and where it is used.
- Describe the k-means algorithm and its objective.
- Explain RL terminology (state, action, reward, policy) at a high level.

Learning Outcomes

By the end of this lecture, you should be able to:

- Explain clustering and where it is used.
- Describe the k-means algorithm and its objective.
- Explain RL terminology (state, action, reward, policy) at a high level.
- Describe semi-supervised learning and when it is helpful.

Abbreviations & Symbols

- Clustering: k clusters, μ_j centroid of cluster j

Abbreviations & Symbols

- Clustering: k clusters, μ_j centroid of cluster j
- Objective (inertia): $\sum_i \|x_i - \mu_{c(i)}\|^2$

Abbreviations & Symbols

- Clustering: k clusters, μ_j centroid of cluster j
- Objective (inertia): $\sum_i \|x_i - \mu_{c(i)}\|^2$
- RL: s state, a action, r reward, π policy, γ discount

Abbreviations & Symbols

- Clustering: k clusters, μ_j centroid of cluster j
- Objective (inertia): $\sum_i \|x_i - \mu_{c(i)}\|^2$
- RL: s state, a action, r reward, π policy, γ discount
- $Q(s, a)$: action-value (Q-function)

Abbreviations & Symbols

- Clustering: k clusters, μ_j centroid of cluster j
- Objective (inertia): $\sum_i \|x_i - \mu_{c(i)}\|^2$
- RL: s state, a action, r reward, π policy, γ discount
- $Q(s, a)$: action-value (Q-function)
- Semi-supervised: labeled set L, unlabeled set U, pseudo-labels

What is Clustering?

Clustering

Group data points into clusters such that points in the same cluster are more similar.

Examples:

- customer segmentation

What is Clustering?

Clustering

Group data points into clusters such that points in the same cluster are more similar.

Examples:

- customer segmentation
- document grouping

What is Clustering?

Clustering

Group data points into clusters such that points in the same cluster are more similar.

Examples:

- customer segmentation
- document grouping
- anomaly detection (as "far from clusters")

k-means (Algorithm Sketch)

1. Choose k cluster centers (initialize).

k-means (Algorithm Sketch)

1. Choose k cluster centers (initialize).
2. Assign each point to the nearest center.

k-means (Algorithm Sketch)

1. Choose k cluster centers (initialize).
2. Assign each point to the nearest center.
3. Update each center as the mean of its assigned points.

k-means (Algorithm Sketch)

1. Choose k cluster centers (initialize).
2. Assign each point to the nearest center.
3. Update each center as the mean of its assigned points.
4. Repeat until convergence.

k-means Objective

k-means minimizes within-cluster squared distance:

$$\sum_{i=1}^{n} \left\| x_i - \mu_{c(i)} \right\|^2$$

where $\mu_{c(i)}$ is the center of the cluster assigned to x_i.

Mini Worked Example (k-means intuition)

Points in 1D: $[1, 2, 8, 9]$, choose $k = 2$.

- Initialize centers: $\mu_1 = 2$, $\mu_2 = 9$

Repeat until assignments stop changing.

Mini Worked Example (k-means intuition)

Points in 1D: $[1, 2, 8, 9]$, choose $k = 2$.

- Initialize centers: $\mu_1 = 2$, $\mu_2 = 9$
- Assign: $\{1, 2\} \to \mu_1$ and $\{8, 9\} \to \mu_2$

Repeat until assignments stop changing.

Mini Worked Example (k-means intuition)

Points in 1D: $[1, 2, 8, 9]$, choose $k = 2$.

- Initialize centers: $\mu_1 = 2$, $\mu_2 = 9$
- Assign: $\{1, 2\} \to \mu_1$ and $\{8, 9\} \to \mu_2$
- Update: $\mu_1 = \text{mean}(1, 2) = 1.5$, $\mu_2 = \text{mean}(8, 9) = 8.5$

Repeat until assignments stop changing.

Choosing k (Practical)

- Elbow method (plot objective vs k)

Choosing k (Practical)

- Elbow method (plot objective vs k)
- Domain knowledge

Choosing k (Practical)

- Elbow method (plot objective vs k)
- Domain knowledge
- Stability across runs (k-means depends on initialization)

Clustering Evaluation (Intuition)

- Internal metrics (no labels): inertia (k-means objective), silhouette score

Clustering Evaluation (Intuition)

- Internal metrics (no labels): inertia (k-means objective), silhouette score
- Visual check (2D/3D after PCA) can reveal separation

Clustering Evaluation (Intuition)

- Internal metrics (no labels): inertia (k-means objective), silhouette score
- Visual check (2D/3D after PCA) can reveal separation
- Beware: "best" clustering depends on the business goal

Reinforcement Learning (High Level)

- Agent interacts with environment: **state** $\rightarrow$ **action** $\rightarrow$ **reward**

Reinforcement Learning (High Level)

- Agent interacts with environment: **state** $\rightarrow$ **action** $\rightarrow$ **reward**
- Goal: learn a **policy** that maximizes cumulative reward

Reinforcement Learning (High Level)

- Agent interacts with environment: **state** → **action** → **reward**
- Goal: learn a **policy** that maximizes cumulative reward
- Feedback is delayed and noisy (harder than supervised learning)

Return and Discounting (Key RL Idea)

- Goal is not just immediate reward, but **total return**:

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$$

Return and Discounting (Key RL Idea)

- Goal is not just immediate reward, but **total return**:

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$$

- $\gamma \in [0, 1)$ discounts future rewards.

Return and Discounting (Key RL Idea)

- Goal is not just immediate reward, but **total return**:

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$$

- $\gamma \in [0, 1)$ discounts future rewards.
- Higher γ values emphasize long-term planning.

Q-Learning (One-line Update)

- Learn an action-value function $Q(s, a)$.

Q-Learning (One-line Update)

- Learn an action-value function $Q(s, a)$.
- Update rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha\big(r + \gamma \max_{a'} Q(s', a') - Q(s, a)\big)$$

Q-Learning (One-line Update)

- Learn an action-value function $Q(s, a)$.
- Update rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha\big(r + \gamma \max_{a'} Q(s', a') - Q(s, a)\big)$$

- α is learning rate; s' is next state after taking action a in state s.

Toy Example: Grid World

- State: agent position on a grid

Toy Example: Grid World

- State: agent position on a grid
- Actions: up/down/left/right

Toy Example: Grid World

- State: agent position on a grid
- Actions: up/down/left/right
- Reward: +1 at goal, -1 at trap, small penalty per step

Toy Example: Grid World

- State: agent position on a grid
- Actions: up/down/left/right
- Reward: +1 at goal, -1 at trap, small penalty per step
- Learn a policy to reach goal quickly

Semi-Supervised Learning (Idea)

- Use a small labeled dataset + large unlabeled dataset.

Risks: error reinforcement if pseudo-labels are wrong.

Semi-Supervised Learning (Idea)

- Use a small labeled dataset + large unlabeled dataset.
- Useful when labeling is expensive but unlabeled data is abundant.

Risks: error reinforcement if pseudo-labels are wrong.

Semi-Supervised Learning (Idea)

- Use a small labeled dataset + large unlabeled dataset.
- Useful when labeling is expensive but unlabeled data is abundant.
- Example approach: **self-training** (pseudo-label confident predictions).

Risks: error reinforcement if pseudo-labels are wrong.

Semi-Supervised Methods (Examples)

- **Self-training**: add confident pseudo-labels iteratively

Use validation to avoid reinforcing mistakes.

Semi-Supervised Methods (Examples)

- **Self-training**: add confident pseudo-labels iteratively
- **Co-training**: two "views" of features teach each other (when applicable)

Use validation to avoid reinforcing mistakes.

Semi-Supervised Methods (Examples)

- **Self-training**: add confident pseudo-labels iteratively
- **Co-training**: two "views" of features teach each other (when applicable)
- **Graph-based label propagation**: spread labels to neighbors

Use validation to avoid reinforcing mistakes.

Summary

- Clustering finds structure without labels; k-means is a common baseline.
- RL learns from reward signals in interaction.
- Semi-supervised learning combines labeled and unlabeled data.

Sources: Mueller and Massaron, *Machine Learning for Dummies*, For Dummies, 2016.;James, Witten, Hastie, and Tibshirani, *An Introduction to Statistical Learning* (2e), selected chapters.;Scikit-learn documentation, accessed 2026-02-28.

Exit Questions

1. What is the k-means objective function (in words)?
2. Why can k-means give different clusters on different runs?
3. Define state, action, reward, and policy in RL.
4. Give one benefit and one risk of semi-supervised learning.